

# Introduction to Machine Learning

Dr. Víctor Uc Cetina

Universität Hamburg

# Content

- 1 Machine learning
- 2 Regarding this ML course
- 3 Supervised Learning
- 4 Unsupervised Learning
- 5 Reinforcement Learning

# What is Machine Learning?

- Machine learning (ML) refers to algorithms used to extract patterns from data and learn a mathematical model that could be used by a computer program to make intelligent decisions.

## Some Data

```
010101010100  
010010101010  
101000101010  
100101001100  
100101000111  
0...
```

## A Model

$$y = f(x)$$

## Decision Making

```
if y=a then do A  
else if y=b then do B  
else if y=c then do C  
...
```


# Supervised learning - Regression

- Given the characteristics and prices of several houses, develop a software for predicting the price of new houses.

Living area	Bedrooms	Price
560	2	37.000
1012	3	79.000
893	3	76.000
2196	4	130.000
$\vdots$	$\vdots$	$\vdots$
936	3	72.000

# Supervised learning - Classification

- Given a set of digital images of blood samples containing Chagas parasites, decide if a digital image of a new blood sample contains at least a Chagas parasite.



# Unsupervised learning - Clustering

- Given a set of text phrases written in two different languages, decide if a new text phrase belongs to one of the languages existing in your set of phrases.

hola a todo el mundo - no me gusta decir  
adiós - hello people - adoro la comida - our  
world is wonderful - el planeta agua -  
ciencia ficción es ciencia - el algoritmo más  
rápido - la mesa es redonda - the door is  
black -  
my chair is broken - el plato está limpio,  
esa escalera está muy inclinada - my  
mouse is wireless - an electronic book - a  
wide road is better, we were at home - ...

# Reinforcement learning - Control

- Given a history of the commands used to control a drone, decide which is the best command to perform in order to avoid a collision with the ground.



# Three major kinds of ML problems

Based on the type of data available and the decisions needed, we can talk about three general kinds of machine learning algorithms:

- Supervised learning: **inputs and outputs**
- Unsupervised learning: **only inputs**
- Reinforcement learning: **states, actions, and rewards**



# Fields of science getting involved in ML


Current machine learning developments come from fields such as:

- Computer Science
- Artificial Intelligence
- Bioinformatics
- Neuroscience
- Psychology
- Robotics


# Practical aspects of machine learning

- Using machine learning algorithms usually involves at least two phases:

## LEARNING PHASE



## APPLICATION PHASE





# Learning machine learning

Recipe for an optimal learning of machine learning:

- ① Understanding the kind of machine learning problem (SL, UL, RL)
- ② Understanding the mathematics (basic statistics, probability, calculus, linear algebra)
- ③ Understanding the algorithms (data structures, complexity)
- ④ Coding (Matlab/Octave/Python)
- ⑤ Experimenting (running programs and plotting graphs)

# Part I - Supervised Learning

- ① Linear regression
- ② Classification
- ③ Generative learning models
- ④ Support vector machines 
- ⑤ AdaBoost 
- ⑥ Learning theory
- ⑦ Regularization and model selection
- ⑧ Neural networks

# Part II - Unsupervised Learning

- ① Expectation-Maximization (EM) algorithm
- ② Principal component analysis
- ③ Self organizing maps
- ④ Spectral clustering

# Part III - Reinforcement Learning

- ① Markov decision processes
- ② Dynamic programming
- ③ Monte Carlo methods
- ④ Temporal difference learning

# Sources of Information

- **Presentations** (Videos and slides)
- **Lecture notes**
- **Exercises**
- **Pattern Recognition and Machine Learning.** Christopher M. Bishop. Springer. 2006.
- **Reinforcement Learning: An Introduction (second edition).** Richard Sutton and Andrew Barto. The MIT Press. 2018.
- **Foundations of Machine Learning.** Mehryar Mohri, Afshin Rostamizadeh and Ameet Talwalkar. The MIT Press. 2012.
- **Understanding Machine Learning: from Theory to Algorithms.** Shai Shalev-Shwartz and Shai Ben-David. Cambridge University Press. 2014.

# Regression

- In the **regression** problem

we have

$$R : X \rightarrow Y,$$


$$X \subseteq \mathbb{R}^d,$$

$$Y \subseteq \mathbb{R}. \text{ (Continuous)}$$

Example:

$$y =$$

$$w_0 + w_1x + w_2x^2 + w_3x^3$$






# Classification

- In the **classification** problem we have  
 $C : X \rightarrow Y$ ,  
 $X \subseteq \mathbb{R}^d$ ,  
 $Y = \{y_1, y_2, \dots, y_n\}$ .  
(Discrete)  
Example:

$$y = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$



# Clustering

The unsupervised learning problem is similar to the classification one, but the data is not labeled. In this case we do not know the category of each example.

- In the **clustering** problem we have


$$C : X \rightarrow T,$$

$$X \subseteq \mathbb{R}^d,$$

$$T = \{t_1, t_2, \dots, t_n\},$$


with unknown  $n$ . Example:

$$t = \arg \min_{t_i} \text{distance}(t_i, x)$$



# Reinforcement Learning

In reinforcement learning an agent (i.e. a robot) must learn to perform a task: a sequence of actions to go from one **start** state to one **goal** state.





Markov decision process  
 $\text{MDP} = (S, A, T, R)$

- Set of states  
 $S = \{s_1, s_2, \dots\}$
- Set of actions  
 $A = \{a_1, a_2, \dots\}$
- Transition function  
(dynamics of the environment)  
 $T : S \times A \rightarrow S$
- Reward function  
 $R : S \rightarrow \mathbb{R}$

# Possible RL solution to the maze problem

We discretize states and actions: 64 states and 4 actions.



**Reward Function:**

$r = -1$  if state is not terminal

$r = 10$  if state is terminal

The agent interacts with the environment several times and the RL algorithm estimates a control function  $a_{t+1} = \pi(s_t)$ .

# Function $Q(s,a)$

S	A	$Q(s,a)$
1	1	0
1	2	0
1	3	0
1	4	0
2	1	0
2	2	0
2	3	0
...	...	...
64	4	0

S	A	$Q(s,a)$
1	1	-20
1	2	-23
1	3	-5
1	4	-22
2	1	-20
2	2	-5
2	3	-22
...	...	...
64	4	-24

# Q-learning Algorithm

Initialize  $Q(s, a)$  arbitrarily

**Repeat** (for each episode):

Initialize  $s$

**Repeat** (for each step of episode):

Choose  $a$  from  $s$  using policy derived from  $Q$  ( $\epsilon$ -greedy)

Take action  $a$ , observe  $r, s'$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$$s \leftarrow s'$$

Until  $s$  is terminal

**Thank you!**

Dr. Víctor Uc Cetina  
cetina@informatik.uni-hamburg.de